

FYP JC2104

Presentation

Imsong Jeon 1155087995
8 December 2021

1. Objectives
2. Process
3. Results
4. Plans and Ideas
5. Conclusion

1. Objectives

FYP JC2104: Open Topics on Distributed Systems for Data Analytics

1. Self-Learning
2. Defining Problem and Solution

2. Process

2021-2022 Term 1

Summer

Beginning

Middle

End

2. Process

Summer



Beginning

Ursa

C++

CMake

Cluster and Job Configuration

2. Process

Middle

Example Codes

hello_world.cc

pi.cc

word_count.cc

kmeans.cc

...

from close point of view:

- High-level APIs
- Source Code

from far point of view:

- Workflow
- Structure

2. Process

Middle

from close point of view:

- High-level APIs
- Source Code

```
100 auto data = TextSourceDataset(input, tg, n_partitions)
101 .FlatMap([features])(const std::string& line) { return LibsvmParse(line, features); })
102 .PartitionBy(
103     [n_partitions](const DataObj& vec) {
104         std::random_device rd;
105         std::mt19937 mt(rd());
106         std::uniform_int_distribution<int> dist(0, n_partitions);
107         return dist(mt);
108     },
109     n_partitions);
```

```
23
24 #include "common/engine.h"
25
```

kmeans.cc

```
17 #include "base/properties.h"
18 #include "common/closure.h"
19 #include "common/dataset/dataset_partition.h"
20 #include "common/dataset/source_dataset.h"
21 #include "common/dataset/table_dataset.h"
22 #include "common/job_driver.h"
23 #include "common/resource_predictor.h"
24 #include "common/task_graph.h"
25
```

common/engine.h

```
20
21 #include "common/constants.h"
22 #include "common/dataset/dataset.h"
23 #include "common/dataset/dataset_partition.h"
24 #include "common/io/input/csv_line_inputformat.h"
25 #include "common/io/input/hdfs_file_splitter.h"
26 #include "common/io/input/hdfs_input_block_info.h"
27 #include "common/io/input/line_inputformat.h"
28 #include "common/io/input/nfs_file_splitter.h"
29 #include "common/io/input/nfs_input_block_info.h"
30 #include "common/source_data.h"
31
```

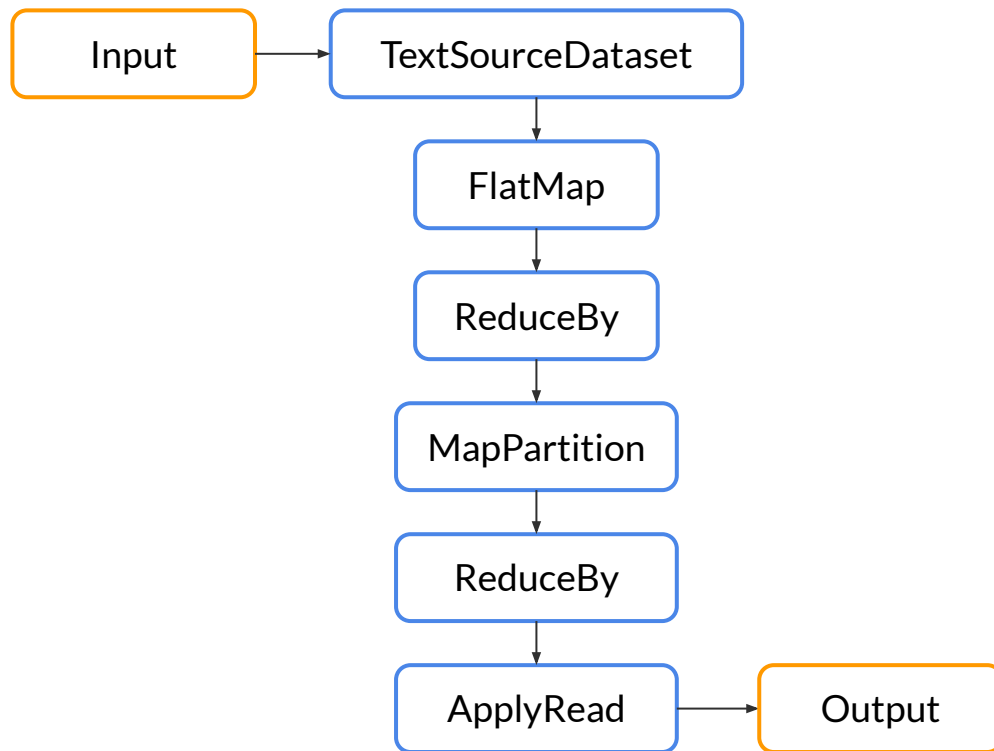
common/dataset/source_dataset.h

2. Process

Middle

from far point of view:

- Workflow (DAG)
- Overall Structure



Workflow of word_count.cc

2. Process

End

More Example Codes

Implementations

Word Length Count

Parallel Dijkstra

3. Results

Word Length Count

Input:

“I have two cats and two dogs”

Word Count output:

```
I1207 00:28:26.212761 182415 word_count.cc:62] 6
```

Word Length Count output:

```
I1207 00:29:46.871788 183663 word_length_count.cc:59] output: 1: 1  
I1207 00:29:46.871800 183663 word_length_count.cc:59] output: 3: 3  
I1207 00:29:46.871802 183663 word_length_count.cc:59] output: 4: 3
```

3. Results

Word Length Count/Run

```
42 class WordCountJob : public Job {
43     public:
44     void Run(TaskGraph* tg, const std::shared_ptr<Properties>& config) const override {
45         TextSourceDataset(config->Get("input"), tg, std::stoi(config->Get("parallelism")))
46         //load data
47         .FlatMap([](const std::string& line) {
48             DatasetPartition<std::pair<int, int>> ret;
49             ParseLine(ret, line);
50             return ret; //dataset of (length, number of word length in this line)
51         })
52         .ReduceBy([](const std::pair<int, int>& ele) { return ele.first; },
53             [](std::pair<int, int>& agg, const std::pair<int, int>& update) { agg.second += update.second; }, 1)
54         .ApplyRead([](auto data) {
55             int count;
56             count = data.size();
57             for(int i = 0; i < count; i++){
58                 std::pair<int, int> cur_data = data.at(i);
59                 LOG(INFO) << "output: " << cur_data.first << ": " << cur_data.second;
60             }
61             google::FlushLogFiles(google::INFO);
62         });
63     }
64 };
```

3. Results

Word Length Count/ParseLine


```
27 void ParseLine(DatasetPartition<std::pair<int, int>>& collection, const std::string& line) {
28     if (line.empty()) {
29         return;
30     }
31     axe::base::WhiteSpaceTokenizer tokenizer(line);
32     std::string tok;
33     std::unordered_map<int, int> length_count;
34     while (tokenizer.next(tok)) {
35         length_count[tok.length()] += 1;
36     }
37     for (auto& pair : length_count) {
38         collection.push_back(pair);
39     }
40 }
41
```

3. Results

Parallel Dijkstra

Input:

1	1	2	7
2	1	3	20
3	2	3	3
4	3	1	5
5	4	1	9
6	5	6	10



Configuration:

n-iterations: 2

start node: 1

Output:

```
I1207 00:48:46.596313 122074 parallel_dij_imsong.cc:225] NodeID = 1 Distance = 0 Parent NodeID = -1
I1207 00:48:46.596325 122074 parallel_dij_imsong.cc:225] NodeID = 3 Distance = 10 Parent NodeID = 2
I1207 00:48:46.596328 122074 parallel_dij_imsong.cc:225] NodeID = 2 Distance = 7 Parent NodeID = 1
```

3. Results

Parallel Dijkstra/Node class

```
28 class Node{
29 public:
30     Node() : neighbor_list_(std::make_shared<std::vector<std::pair<int,int>>>()) {}
31     Node(int node_id, int min_distance, int parent_node_id) :
32         node_id_(node_id), min_distance_(min_distance), parent_node_id_(parent_node_id),
33         neighbor_list_(std::make_shared<std::vector<std::pair<int,int>>>()) {}
34     Node(int node_id, int min_distance, int parent_node_id, const std::shared_ptr<std::vector<std::pair<int,int>>>& neighbor_list) :
35         node_id_(node_id), min_distance_(min_distance), parent_node_id_(parent_node_id), neighbor_list_(neighbor_list) {}
36
37     int GetNodeID() const { return node_id_; }
38     int GetMinDistance() const { return min_distance_; }
39     int GetParentNodeID() const { return parent_node_id_; }
40     const std::shared_ptr<std::vector<std::pair<int,int>>>& GetNeighborList() const { return neighbor_list_; }
41
42     double GetMemory() const {
43         double ret = sizeof(int) * 3 + neighbor_list_>size() * sizeof(std::pair<int,int>);
44         return ret;
45     }
46     bool operator<(const Node& other) const { return node_id_ < other.node_id_; }
47     bool operator==(const Node& other) const { return node_id_ == other.node_id_; }
48
49     friend void operator<<(axe::base::BinStream& bin_stream, const Node& n) {
50         bin_stream << n.node_id_ << n.min_distance_ << n.parent_node_id_ << *(n.neighbor_list_); }
51     friend void operator>>(axe::base::BinStream& bin_stream, Node& n) {
52         bin_stream >> n.node_id_ >> n.min_distance_ >> n.parent_node_id_ >> *(n.neighbor_list_); }
53
54 private:
55     int node_id_;
56     int min_distance_;
57     int parent_node_id_;
58     std::shared_ptr<std::vector<std::pair<int,int>>> neighbor_list_;
59 };
60
```

3. Results

Parallel Dijkstra/Run (part 1)

```
110 class ParallelDijkstra : public Job {
111     public:
112     void Run(TaskGraph* tg, const std::shared_ptr<Properties>& config) const override {
113         auto input = config->GetOrSet("data", "");
114         int n_partitions = std::stoi(config->GetOrSet("parallelism", "20"));
115         int n_iters = std::stoi(config->GetOrSet("n_iters", "3"));
116         int start_ID = std::stoi(config->GetOrSet("start_node", "1"));
117         // Load data
118         auto graph = TextSourceDataset(input, tg, n_partitions)
119             .FlatMap([start_ID](const std::string& line) { return ParseLine(line, start_ID); })
120             .ReduceBy([](const Node& ele) { return ele.GetNodeID(); },
121                 [](Node& agg, const Node& ele) {
122                     std::shared_ptr<std::vector<std::pair<int,int>>> aggList = agg.GetNeighborList();
123                     std::shared_ptr<std::vector<std::pair<int,int>>> eleList = ele.GetNeighborList();
124                     aggList->push_back(eleList->front());
125                 },
126                 n_partitions);
127     }
```

I1207 11:46:07.515542 256070 parallel_dij_umsong.cc:104] Node ID = 1 minDist = 0 parent_node_id = -1 neighbors = 2 7 3 20

3. Results

Parallel Dijkstra/Run (part 2)

```
172 auto updated_graph = std::make_shared<axe::common::Dataset<Node>>(graph);
173 std::shared_ptr<axe::common::Dataset<std::pair<int, std::pair<int, int>>>> min_distances;
174
175 for(int i = 0; i < n_iters; ++i){
176     min_distances = std::make_shared<axe::common::Dataset<std::pair<int, std::pair<int, int>>>>(updated_graph->MapPartition(get_distances)
177     .ReduceBy([](const std::pair<int, std::pair<int, int>>& ele) { return ele.first; },
178     [](std::pair<int, std::pair<int, int>>& agg, const std::pair<int, std::pair<int, int>>& ele) {
179         if(agg.second.first > ele.second.first){
180             agg.second.first = ele.second.first; //update min distance
181             agg.second.second = ele.second.second; //update parent ID
182         }
183     },
184     n_partitions));
185
186     updated_graph = std::make_shared<axe::common::Dataset<Node>>{
187         updated_graph->SharedDataMapPartitionWith(min_distances.get(), apply_updates));
188 }
```

```
139 auto get_distances = [](const DatasetPartition<Node>& graph){
140     DatasetPartition<std::pair<int, std::pair<int, int>>> distances;
141     for(const Node& n : graph){
142         distances.push_back(std::make_pair(n.GetNodeID(), std::make_pair(n.GetMinDistance(), n.GetParentNodeID())));
143         for(std::pair<int,int> neighbor : *n.GetNeighborList()){
144             int new_distance = MAX_INT;
145             if(n.GetMinDistance() < MAX_INT)
146                 new_distance = neighbor.second + n.GetMinDistance();
147             distances.push_back(std::make_pair(neighbor.first, std::make_pair(new_distance, n.GetNodeID())));
148         }
149     }
150     return distances;
151 };
```


3. Results

Parallel Dijkstra/Run (part 3)

```
153 auto apply_updates = [](const DatasetPartition<Node>& graph, const DatasetPartition<std::pair<int, std::pair<int, int>>>& min_distances) {
154     DatasetPartition<Node> updated_graph;
155     for(const std::pair<int, std::pair<int, int>>& p : min_distances){
156         int node_id = p.first;
157         bool updated = false;
158         for(const Node& n : graph){
159             if(n.GetNodeID() == node_id){
160                 updated_graph.push_back(Node(node_id, p.second.first, p.second.second, n.GetNeighborList()));
161                 updated = true;
162                 break;
163             }
164         }
165         if(!updated){
166             updated_graph.push_back(Node(node_id, p.second.first, p.second.second));
167         }
168     }
169     return updated_graph;
170 };
```

```
197 auto get_output = [](const DatasetPartition<Node>& graph){
198     DatasetPartition<std::pair<int, std::pair<int, int>>> output;
199     for(const Node& n : graph){
200         if(n.GetMinDistance() < MAX_INT){
201             output.push_back(std::make_pair(n.GetNodeID(), std::make_pair(n.GetMinDistance(), n.GetParentNodeID())));
202         }
203     }
204     return output;
205 };
206
207 updated_graph->MapPartition(get_output)
208     .PartitionBy([](const std::pair<int, std::pair<int, int>>&) { return 0; }, 1)
209     .ApplyRead( [](const auto& partition) {
210         for (auto& par : partition) {
211             LOG(INFO) << "NodeID = " << par.first << " Distance = " << par.second.first << " Parent NodeID = " << par.second.second;
212         }
213     });
```

4. Plans and Ideas

Plans:

- Implement and test with larger data and more workers
- More study
- Distributed PCA (Principal Component Analysis)

Ideas

- Linear Algebra Package
- Real-time interaction

5. Conclusion

Thank you!